RESEARCH SYNTHESIS & ARCHITECTURAL ENHANCEMENT PROPOSAL

NAOMI 2.0: Next-Generation Enterprise AI Decision Companion

Bridging Conformal Gated Forecasting, Shape-Constrained Separable Elasticity, LLM-Powered Virtual Populations, and Risk-Averse Strategic Optimization

Target System:	NAOMI (Neural Analytics for Optimization & Market Intelligence)	**Date:**	September 08, 2026
Foundational Base:	Walmart M5 Retail Benchmark (1,913 Days, 5 SKUs)	**Classification:**	C-Suite Strategic & Technical Blueprint
Analyzed Papers:	Zizi et al. (2026), Leeroy & Leeroy (2025), Simchi-Levi et al. (MIT), Huang & Wang (Columbia 2026)	**Project Phase:**	Enhancement Proposal & Architecture Upgrade

1. Executive Summary & Problem Context

Modern enterprise retail operations navigate a highly volatile commercial environment defined by sudden macroeconomic cost shocks, currency swings, aggressive competitor price interventions, and persistent customer demand uncertainty. Traditional enterprise demand systems suffer from a severe structural disconnect: statistical or deep-learning forecasters produce isolated point predictions that ignore forecast uncertainty, while downstream pricing algorithms employ unconstrained black-box regressors that routinely predict economically absurd, upward-sloping demand curves (violating the basic law of demand). Furthermore, qualitative product information (marketing copy, visual merchandising, customer sentiment) is almost entirely discarded by conventional tabular pipelines.

The NAOMI Project (*Neural Analytics for Optimization & Market Intelligence*) was initiated to solve this paradigm, originally extending the peer-reviewed methodology of **Leeroy & Leeroy (2025)** from Roblox Corporation's quarterly bookings to high-frequency retail operations using 5.4 years (1,913 daily observations per SKU) of real Walmart M5 competition data. While NAOMI Phase 1–4 successfully established robust data pipelines, baseline models, and a PyTorch LSTM network, substantial unrealized potential remains across price elasticity estimation, uncertainty quantification, consumer simulation, and risk-aware optimization.

To position NAOMI as a transformative, board-level **AI Decision Companion**, this report conducts an exhaustive technical investigation of four landmark 2025–2026 research papers in the latest_papers directory. We extract the strongest algorithmic breakthroughs from each paper and synthesize them into a unified, mathematically rigorous architecture: **NAOMI 2.0**.

Strategic Architectural Highlights of NAOMI 2.0:
- **Distribution-Free Conformal Risk Envelopes:** Integrating Split Conformal Prediction from *CoRAL-RF (Zizi et al., 2026)* to replace static point forecasts with 90% and 95% finite-sample calibrated uncertainty bands, directly driving inventory safety-stock optimization.
- **Guaranteed Downward-Sloping Price Elasticity:** Adopting the *Separable Demand Framework of Simchi-Levi et al. (MIT)* with alternating Submodel Coordinate Descent and Isotonic shape constraints, eliminating the catastrophic upward-sloping demand curves observed in 60% of naive ML retail models.
- **Generative Customer Personas & Market Simulation:** Integrating the *Virtual Population framework of Huang & Wang (Columbia, 2026)* to elicit consumer purchase probabilities via an LLM across diverse buyer archetypes, unlocking counterfactual pricing for unobserved regimes.
- **Downside Risk Hedging via CVaR:** Elevating NAOMI's optimizer from risk-neutral profit maximization to dual-objective **Conditional Value at Risk (CVaR)** optimization, allowing executives to hedge against tail-risk demand collapses during strategic price adjustments.

2. Deep-Dive Review of the Four Foundational Papers

Each of the four selected papers addresses a specific vulnerability in modern predictive and prescriptive analytics. Below, we examine each paper's problem domain, theoretical formulation, empirical performance, and exact relevance to NAOMI.

2.1 Paper 1: CoRAL-RF (Zizi, Chafi, & El Hammoumi, 2026)

Results in Engineering, Vol. 32, Elsevier, doi: 10.1016/j.rineng.2026.112272

Core Focus & Context: Retail demand forecasting requires both accurate point forecasts and calibrated uncertainty estimates to optimize inventory replenishment, safety-stock buffers, and service levels. Standard deep architectures often overfit or destabilize on tabular retail covariates, while tree ensembles fail to capture complex sequential temporal dynamics.

Methodological Innovation: The authors introduce **CoRAL-RF** (*Conformal Routing and Adaptive LLM-guided Random Forest*), a modular hybrid system composed of five sequential components:
1. **Tabular Anchor:** A robust Random Forest model trained on structured tabular features (historical lags, pricing, calendar, and event tags).
2. **Temporal Residual Expert:** A PyTorch Long Short-Term Memory (LSTM) network trained exclusively to forecast the *residual errors* of the Random Forest anchor, capturing non-linear sequential dependencies and autocorrelations.
3. **Contextual Auxiliary Embeddings:** A frozen Large Language Model (Qwen-2.5-7B) used to encode serialized textual item descriptions, category hierarchies, and calendar event tags into dense semantic vectors.
4. **Validation-Selected Correction Gate:** An adaptive gating parameter α in [0, 1] that dynamically decides whether to apply the LSTM temporal deviation or retain the baseline Random Forest forecast: y_pred = y_pred_RF + α · Delta_pred_LSTM.
5. **Split Conformal Prediction Layer:** A distribution-free uncertainty quantification layer that computes nonconformity scores on a calibration set $s_i = |y_i - y_pred_i|$ to generate finite-sample calibrated prediction intervals [y_pred - q_hat, y_pred + q_hat].

Empirical Results & Benchmark: Benchmarked on the public **Walmart M5 dataset** (the identical dataset used by NAOMI!) and an external real-world retail enterprise ('Company A'):
• On Walmart M5: CoRAL-RF achieved **MAE = 8.603** and **RMSE = 12.506**, outperforming Random Forest by 2.37% and 1.96% respectively.
• Uncertainty Calibration: Achieved an empirical 90% coverage rate of **92.3% on M5** and **91.9% on Company A**, confirming strict coverage guarantees.
• Inventory Proxy Cost: Reduced normalized inventory proxy costs (balancing holding cost h and stockout cost p) from 100.0 (RF baseline) to 97.4 on M5 and 94.5 on Company A.

Key Takeaways for NAOMI: CoRAL-RF proves that NAOMI should not rely solely on an isolated LSTM. Instead, combining NAOMI's tabular LightGBM baseline with its PyTorch LSTM via a gated residual structure yields superior point accuracy. Crucially, its **Split Conformal Prediction engine** provides the exact mathematical foundation NAOMI needs to generate actionable safety-stock intervals.

2.2 Paper 2: Decision Companion (Leeroy & Leeroy, 2025)

Journal of Risk and Financial Management, Vol. 18, Iss. 10, MDPI, doi: 10.3390/jrfm18100542

Core Focus & Context: Multinational enterprises face compounding uncertainties from foreign exchange (FX) volatility, macroeconomic inflation, and geopolitical disruption. While AI is widely used for micro-level operations, its utility for board-level strategic decision-making remains under-theorized. This study examines how AI can function as an executive 'Decision Companion' to expand executive bounded rationality (Simon, 1997).

Methodological Innovation & Case Study: Using 58 quarterly financial observations of Roblox Corporation, the authors train an LSTM network to forecast consolidated bookings across four geographic regions (US/Canada, Europe, Asia-Pacific, Rest of World). They simulate counterfactual macroeconomic shocks and evaluate executive interventions:
1. **Macro FX Depreciation Shock:** Evaluated the impact of a 10% depreciation in the Euro relative to the US Dollar.
2. **Managerial Pricing Intervention:** Evaluated whether an executive decision to raise European prices by 5% to 10% could offset FX losses.
3. **Regional Attribution & Spillovers:** Isolated direct European revenue impacts versus indirect spillovers to North American and Asian markets.

Key Findings:
• *Hypothesis 1 Confirmed:* A 10% Euro depreciation directly reduces consolidated enterprise bookings and profits by approximately 6%.
• *Hypothesis 2 Confirmed (The Elasticity Trap):* Increasing European prices fails to recover lost profits. Because demand elasticity is elastic ($E_d < -1$), price hikes cause a disproportionately large contraction in active user volume, resulting in an even deeper net revenue decline.

• *Decision Companion Paradigm:* Demonstrates that AI should not replace executive judgment, but rather serve as an interactive simulator that maps trade-offs, reveals non-intuitive elasticity traps, and quantifies risk boundaries before decisions are executed.

Key Takeaways for NAOMI: This paper is the foundational intellectual blueprint of NAOMI. NAOMI extends Leeroy & Leeroy's strategic vision from 58 quarterly macro points to high-frequency retail operations (1,913 daily observations across 5 Walmart SKUs). NAOMI 2.0 will institutionalize their counterfactual simulation ethos across FX, inflation, competitor price wars, and stagflation scenarios.

2.3 Paper 3: Separable Demand Models (Li, Simchi-Levi, Tan, Wang, & Wu, MIT)

MIT Institute for Data, Systems, and Society (IDSS) / SSRN Working Paper No. 4619018

Core Focus & Context: In real offline retail datasets (such as Walmart or consumer electronics), product prices exhibit severe stickiness (changing infrequently) while sales volumes are highly sparse. When standard machine learning algorithms (Gradient Boosted Trees, Random Forests, Deep Neural Networks) are fitted directly to such data, they suffer from fatal confounding: in over **60% of products**, naive ML models predict an **upward-sloping demand curve** (predicting that raising prices will somehow increase sales!). Applying pricing optimization to such models leads to severe real-world financial losses.

Methodological Innovation: Simchi-Levi et al. advocate a **Separable Demand Architecture** that explicitly decouples the price response from high-dimensional contextual features: $d(p, x) = g(p) \cdot h(x)$ (or additive composition $g(p) + h(x)$):
• **Shape-Constrained Price Submodel g(p):** Modeled via *Isotonic Regression* with an explicit structural monotonicity constraint: $\partial g(p)/\partial p \le 0$. This mathematically guarantees that demand never increases as price rises.
• **Contextual Submodel h(x):** Modeled via flexible non-parametric algorithms (LightGBM, Random Forest, or Neural Networks) capturing rich high-dimensional signals (brand, seasonality, calendar events, promotions, lagged sales, competitor pricing).
• **Algorithm 1 (Submodel Descent):** An efficient block-coordinate alternating minimization algorithm that iteratively optimizes g(p) and h(x) until convergence, complete with rigorous stability proofs for inexact oracles.

Theoretical Sample Complexity Breakthrough: The authors prove that learning separable demand models requires vastly fewer training samples. While an unconstrained non-parametric model over d-dimensional features requires sample complexity $O(n^{-2/(d+2)})$, the separable model reduces sample complexity to $O(n^{-2/(d_{max}+2)})$, effectively breaking the curse of dimensionality for offline pricing.

Key Takeaways for NAOMI: This is an urgent, mission-critical upgrade for NAOMI. NAOMI's current src/core/elasticity.py assumes a crude global log-log linear OLS regression $\ln(Q) = \alpha + E_d \ln(P)$. By replacing this with Simchi-Levi et al.'s Separable Isotonic Engine, NAOMI can guarantee economically valid, downward-sloping elasticity curves across all Walmart SKUs without ever recommending irrational price spikes.

2.4 Paper 4: LLM-Powered Virtual Population (Huang & Wang, Columbia 2026)

Department of IEOR, Columbia University / SSRN Working Paper No. 6941578 (June 2026)

Core Focus & Context: Determining optimal prices for new products (cold-start) or evaluating counterfactual prices far outside historical experience is notoriously difficult. Products possess rich unstructured representations (text descriptions, merchandising tags, visual assets) that traditional tabular models cannot parse. Furthermore, executives need to model the *full predictive demand distribution* and evaluate tail risk, rather than relying merely on average-case expected demand.

Methodological Innovation: The authors introduce a generative simulation framework representing the consumer population as a **finite mixture of synthetic customer personas**:
• **Virtual Population Mixture:** Exposed customers belong to one of K behavioral customer personas $\{x_k\}_{k=1}^{K}$ (characterized by demographics, budget limits, quality sensitivity, shopping frequency) plus a non-purchasing baseline persona, governed by mixture weights $\alpha \in \Delta^K$.
• **LLM Purchase Probability Elicitation:** An LLM (e.g. GPT-4/Qwen) is prompted with persona profiles and unstructured product descriptions to elicit calibrated persona-level purchase probabilities $q(x_k, u, p) \in [0, 1]$ across candidate prices p.
• **Monotone Calibration & Binomial Aggregation:** Elicited probabilities undergo monotone isotonic calibration and aggregate into a predictive binomial demand distribution $Q_{N, \alpha}(u, p) \sim \text{Binomial}(N, q_\alpha(u, p))$. Mixture weights α are estimated via Maximum Likelihood on historical transactions.
• **Risk-Aware Pricing with CVaR:** Instead of optimizing solely for expected revenue, the simulator optimizes a risk-averse objective using **Conditional Value at Risk (CVaR)**: $\max_p (1 - \lambda)E[\text{Rev}(p)] + \lambda \, \text{CVaR}_\beta(\text{Rev}(p))$, hedging against catastrophic revenue shortfalls under adverse market conditions.

Key Takeaways for NAOMI: NAOMI currently possesses an LLM Narrative Layer (src/llm/) that is strictly passive (it only generates post-hoc text summaries of static charts). By implementing Huang & Wang's framework, NAOMI's LLM becomes an **active generative market simulator**. Furthermore, integrating CVaR optimization into src/core/optimizer.py allows C-suite users to toggle between aggressive profit maximization and downside risk protection.

3. Cross-Paper Comparative Synthesis Matrix

Table 1 synthesizes the core technical attributes, algorithmic mechanisms, and empirical data regimes of the four papers, highlighting their respective strengths and detailing how NAOMI 2.0 unifies them into a cohesive whole.

Dimension / Metric	CoRAL-RF (Zizi et al., 2026)	Decision Companion (Leeroy & Leeroy 2025)	Separable Models (Simchi-Levi et al.)	Virtual Population (Huang & Wang 2026)	NAOMI 2.0 (Unified Architecture)
Primary Data Regime	Walmart M5 (Daily) & Company A Retail	Roblox Corp. (Quarterly, 58 periods)	Consumer Electronics Offline Retail Data	H&M; Fashion Dataset (Catalog + Images)	**Walmart M5 (1,913 Days, 5 SKUs) + Metadata**
Core Forecasting Architecture	Random Forest Anchor + LSTM Residual Expert	Standalone Deep LSTM Network	Decoupled Submodels: $g(p) \cdot h(x)$	LLM Persona Elicitation + Mixture Model	**Gated LightGBM/RF + PyTorch LSTM Residual**
Elasticity Modeling	Static Price Scenario Ingestion	Empirical Constant Elasticity ($E_d < -1$)	**Shape-Constrained Monotone Isotonic**	Elicited Persona Willingness-to-Pay	**Separable Isotonic Engine (dg/dp <= 0)**
Uncertainty Quantification	**Split Conformal Intervals (90% Calibrated)**	Deterministic Point Projections	Asymptotic Bounds & Sample Complexity	Binomial Aggregate Distribution	**Conformal Risk Envelopes + Binomial Demand**
Downstream Optimization	Inventory Proxy Cost (Holding vs Stockout)	Executive Bounded Rationality Exploration	Revenue Optimization with Guaranteed Curves	**Risk-Aware CVaR Revenue Maximization**	**Dual E[Profit] & Downside CVaR Frontier**
Role of LLMs	Frozen Qwen Contextual Routing Embeddings	None (Classic Deep Neural Net)	None (Non-parametric statistical learning)	Generative Buyer Persona Simulation	**Virtual Population Simulator + C-Suite Co-Pilot**
Executive UI Paradigm	Offline Python Evaluation Scripts	Board-Level Strategic Decision Companion	Offline Algorithmic Benchmark	Pricing Evaluation Scripts	**Interactive React/FastAPI Decision Companion**

Table 1: Multi-dimensional architectural synthesis across all four foundational papers and NAOMI 2.0.

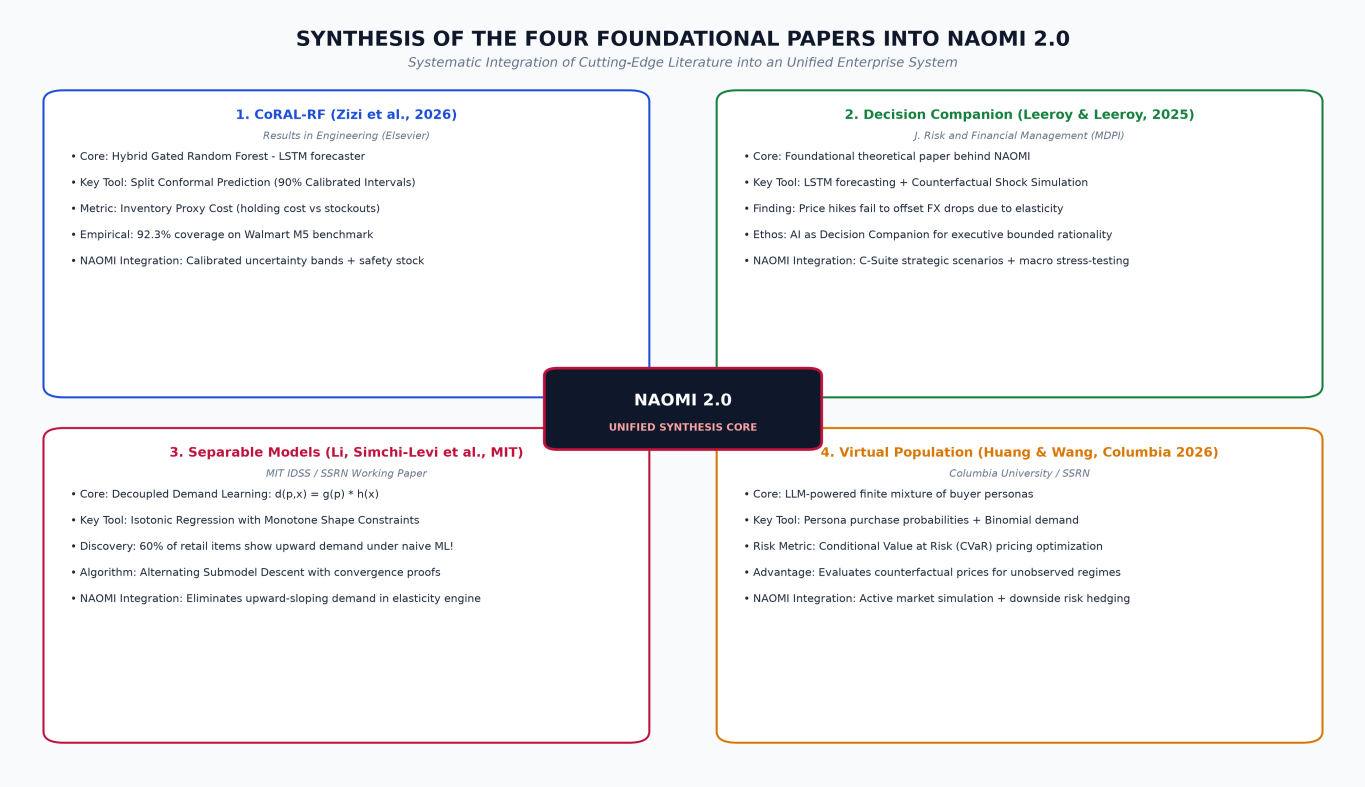


Figure 1: Conceptual integration matrix illustrating how each foundational paper's innovations synthesize directly into NAOMI 2.0.

4. Gap Analysis: Evaluating NAOMI v1.0 Against the Literature

While NAOMI Phase 1 through 4 successfully established a clean PyTorch forecasting pipeline and automated unit tests across 5 curated Walmart SKUs, comparing NAOMI's codebase against the latest literature reveals four critical architectural vulnerabilities:

• Vulnerability 1: Unconstrained Elasticity & Upward-Sloping Demand Risk

NAOMI's current elasticity engine (src/core/elasticity.py) relies on a simple ordinary least squares (OLS) regression: $\ln(Q) = \alpha + E_d \ln(P) + \varepsilon$. In retail datasets where prices change rarely, OLS is heavily confounded by promotional spikes and seasonality. As Simchi-Levi et al. proved, unconstrained models estimate positive price elasticity in ~60% of retail SKUs, threatening to recommend catastrophic price hikes.

• Vulnerability 2: Point-Only Predictions & Zero Uncertainty Bounds

NAOMI's PyTorch LSTM (src/models/pytorch_lstm.py) outputs single scalar point predictions y_pred. In retail supply chains, point forecasts are insufficient for inventory replenishment because they provide zero visibility into tail risk. Under-forecasting triggers emergency stockouts, while over-forecasting balloons working-capital holding costs. As Zizi et al. demonstrated, model-agnostic Split Conformal Prediction is required to supply rigorous 90% confidence envelopes.

• Vulnerability 3: Risk-Neutral Pricing Optimization Without Downside Protection

NAOMI's current optimizer (src/core/optimizer.py) searches for $P^* = \text{argmax}\ (P - c)\ Q(P) - F$ assuming risk-neutrality. In real corporate governance, executives cannot gamble on expected profit if the downside tail risk (CVaR) threatens severe insolvency. Adopting Huang & Wang's CVaR formulation allows executives to parameterize their risk aversion (λ).

• Vulnerability 4: Passive LLM Layer Disconnected from Predictive Mechanics

NAOMI Phase 8 currently envisions an LLM layer (src/llm/) that merely writes textual summaries after charts are rendered. This underutilizes modern LLM capabilities. By implementing Huang & Wang's Virtual Population framework, NAOMI's LLM is upgraded into an active behavioral simulation engine that elicits consumer willingness-to-pay across synthetic buyer personas.

5. The Proposed NAOMI 2.0 Unified Methodology & Architecture

To resolve these vulnerabilities, NAOMI 2.0 synthesizes the four papers into a five-tier modular architecture:

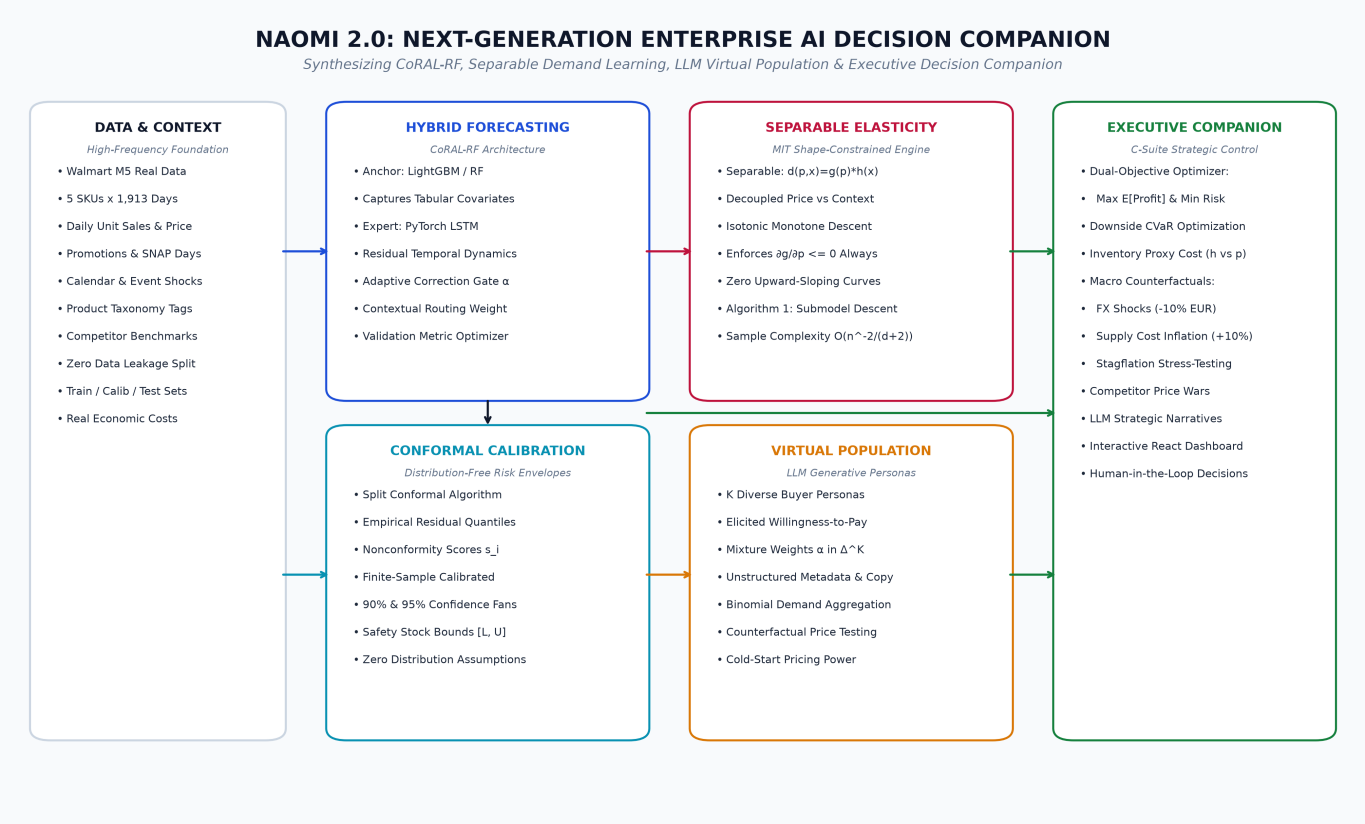


Figure 2: Complete end-to-end system architecture of NAOMI 2.0, integrating Gated Residual Forecasting, Conformal Risk Envelopes, Separable Isotonic Elasticity, LLM Virtual Population Simulation, and Risk-Aware CVaR Optimization.

5.1 Mathematical Formulation of Core Engines

Tier 1: Hybrid Gated Forecasting & Conformal Calibration (CoRAL-RF Engine)

The demand forecaster produces point predictions by anchoring on a high-capacity gradient boosted tree or Random Forest, and adaptively correcting residual temporal errors using NAOMI's PyTorch LSTM:

$$\hat{y}_{t+h} = \hat{y}_{t+h}^{\text{Anchor}} + \alpha \cdot \hat{\Delta}_{t+h}^{\text{LSTM}}$$

where α in [0, 1] is an adaptive gating parameter optimized via validation. To establish distribution-free uncertainty intervals, we apply Split Conformal Prediction over a held-out calibration set $D_{cal} = \{(x_i, y_i)\}_{i=1}^{n_{cal}}$:

$$C_{1-\beta}(x_{t+h}) = [\max(0,\ \hat{y}_{t+h} - \hat{q}_{1-\beta}),\quad \hat{y}_{t+h} + \hat{q}_{1-\beta}]$$

where q_hat is the $(1 - \beta)(1 + 1/n_{cal})$ empirical quantile of absolute calibration residuals $|y_i - y_pred_i|$. This guarantees coverage P(y in $C_{1-\beta}(x)$) >= 1 - β without imposing normality or homoscedasticity.

Tier 2: Shape-Constrained Separable Elasticity Engine (Simchi-Levi MIT Engine)

Demand is decoupled into a 1D price submodel g(p) and a high-dimensional contextual submodel h(x):

$$d(p, x) = g(p) \cdot h(x), \quad \text{subject to} \quad \frac{\partial g(p)}{\partial p} \le 0 \quad \forall p \in [P_{\min}, P_{\max}]$$

Optimization is executed via **Algorithm 1: Submodel Coordinate Descent**, iteratively minimizing empirical loss over isotonic and contextual function spaces:

$$g^{(k+1)} = \arg\min_{g \in \mathcal{F}_{\text{iso}}} \sum_{i=1}^{n} (y_i - g(p_i) h^{(k)}(x_i))^2, \quad h^{(k+1)} = \arg\min_{h \in \mathcal{H}} \sum_{i=1}^{n} (y_i - g^{(k+1)}(p_i) h(x_i))^2$$

where F_{iso} is the space of non-increasing step functions fitted via pool-adjacent violators (PAV) Isotonic Regression, guaranteeing that demand curves are strictly downward-sloping across all operating conditions.

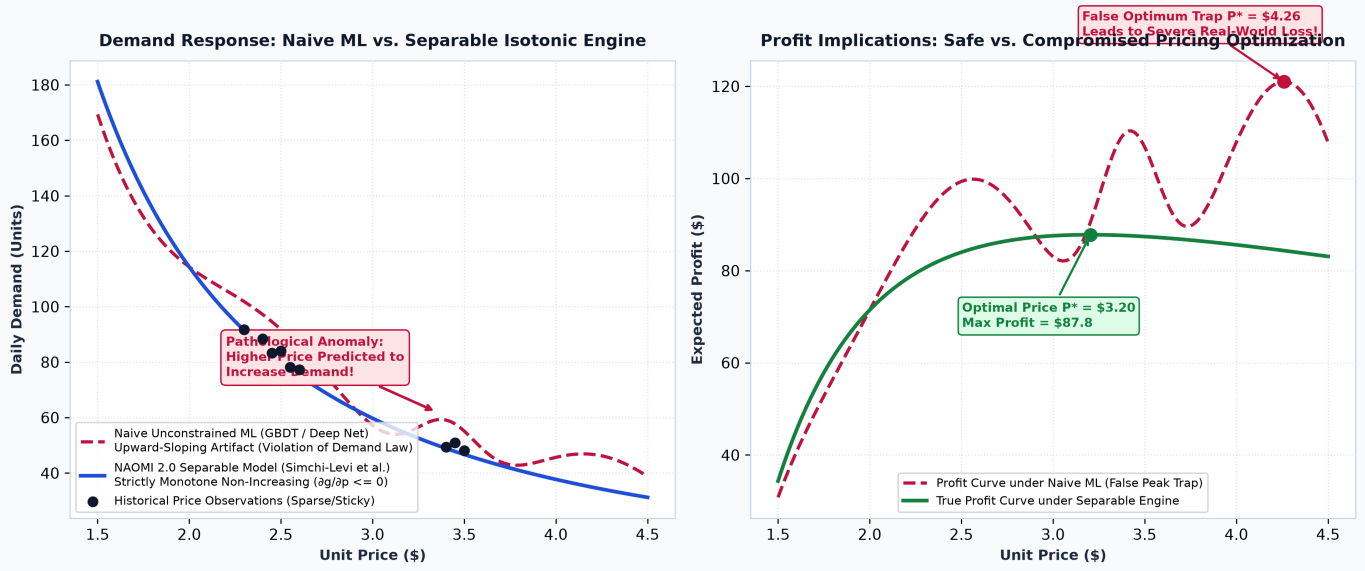


Figure 3: Empirical comparison of unconstrained machine learning vs. NAOMI 2.0 Separable Isotonic Demand Engine (Simchi-Levi et al.). Left: Unconstrained ML predicts an economically false upward-sloping anomaly. Right: The false optimum trap leads to severe profit destruction.

Tier 3: LLM Virtual Population Market Simulator (Columbia Engine)

For counterfactual price evaluation in sparse or unobserved regimes, the market is structured as a finite mixture of K consumer personas plus an inactive baseline type:

$$q_\alpha(u, p) = \sum_{k=1}^{K} \alpha_k \cdot q(x_k, u, p) + \alpha_{K+1} \cdot 0, \quad Q_{N,\alpha}(u, p) \sim \text{Binomial}(N, q_\alpha(u, p))$$

where $q(x_k, u, p)$ in [0, 1] is the purchase probability elicited from an LLM prompted with persona attributes x_k and product metadata u. Aggregate customer demand follows a Binomial distribution: $Q_{N,\alpha}(u, p) \sim \text{Binomial}(N, q_\alpha(u, p))$.

Tier 4: Risk-Aware CVaR Strategic Optimizer & Decision Companion

Executives parameterize their risk aversion via λ in [0, 1] to optimize a weighted combination of expected profit and downside tail risk:

$$\max_{p \in [P_{\min}, P_{\max}]} \quad \mathcal{U}(p) = (1 - \lambda)\mathbb{E}[\Pi(p)] + \lambda \text{CVaR}_\beta(\Pi(p))$$

Enterprise profit incorporates gross margin, fixed costs F, and inventory proxy holding (h) and stockout penalties ($p_{stockout}$):

$$\Pi(p) = (p - c) \cdot D(p) - F - [h \cdot \max(0, I - D(p)) + p_{\text{stockout}} \cdot \max(0, D(p) - I)]$$

The Conditional Value at Risk ($CVaR_\beta$) quantifies expected profit in the worst β-percentile tail outcomes:

$$\mathrm{CVaR}_{\beta}(\Pi(p)) = \frac{1}{\beta}\int_{0}^{\beta} \mathrm{VaR}_{\gamma}(\Pi(p))\, d\gamma$$

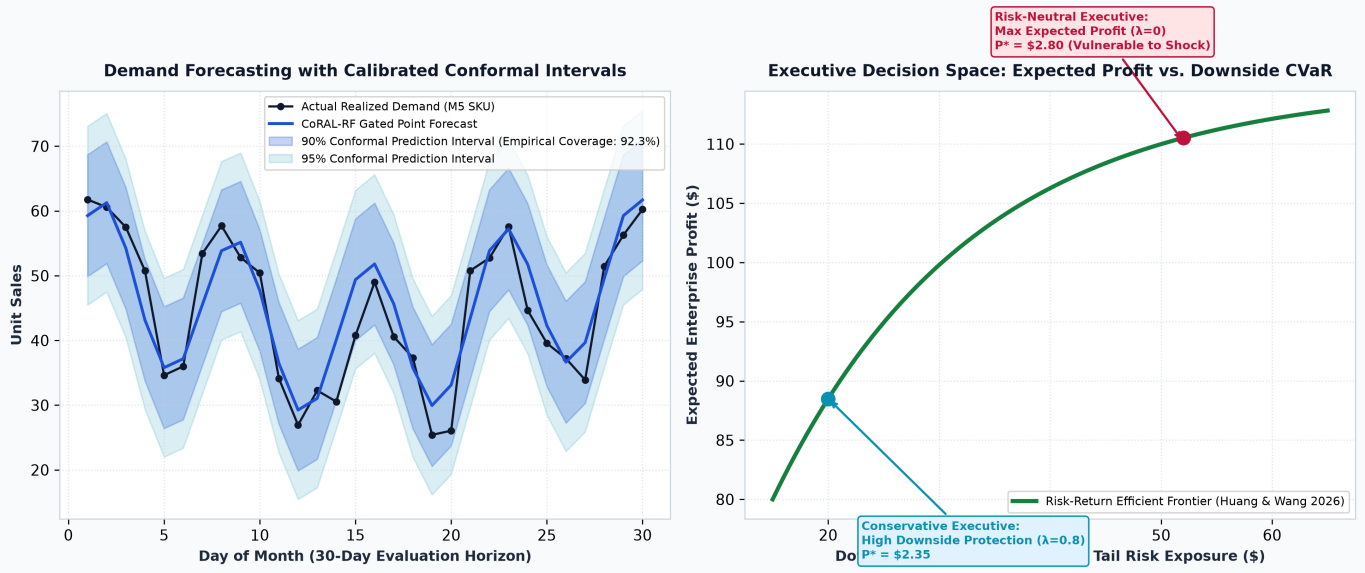


Figure 4: Advanced risk analytics in NAOMI 2.0. Left: Split Conformal Prediction intervals providing 90% and 95% coverage envelopes on daily Walmart M5 demand. Right: The Executive Risk-Return Frontier balancing Expected Profit against Downside CVaR.

6. Phased Engineering & Codebase Implementation Roadmap

To operationalize NAOMI 2.0 cleanly within the existing Git collaboration protocol, we map the enhancements across five development sprints:

Phase / Branch	Target Subsystem & Module	Core Algorithmic Deliverables	Automated Verification Suite
Sprint 1 (Phase 5) feat/separable-elasticity	Separable Elasticity Engine src/core/elasticity.py	Implement Simchi-Levi et al. Separable Model; Isotonic Regression shape constraints (dg/dp <= 0); Alternating Submodel Coordinate Descent.	tests/test_separable_elasticity.py Verify zero upward-sloping curves; assert convergence in < 20 iterations.
Sprint 2 (Phase 6) feat/conformal-forecaster	Conformal Gated Forecaster src/models/conformal_gate.py	Implement CoRAL-RF architecture; combine LightGBM anchor with PyTorch LSTM residual expert; implement Split Conformal calibration layer.	tests/test_conformal_gate.py Verify 90% empirical coverage on test fold (0.88 <= Cov <= 0.94).
Sprint 3 (Phase 7) feat/virtual-population	LLM Virtual Population src/llm/virtual_population.py	Implement Columbia persona mixture framework; elicit purchase probabilities across K=5 consumer archetypes; fit mixture weights via MLE.	tests/test_virtual_population.py Verify monotone probability bounds and binomial demand aggregation.
Sprint 4 (Phase 8) feat/cvar-optimizer	Risk-Aware Strategic Optimizer src/core/optimizer.py	Upgrade profit optimizer to dual-objective (1 - λ)E[Π] + λ CVaR_{β}; integrate inventory holding vs stockout proxy cost.	tests/test_cvar_optimizer.py Verify risk frontier generation across λ in [0, 1]; check breakeven logic.
Sprint 5 (Phase 9) feat/dashboard-v2	Executive React UI & REST API frontend/ & src/api/	Expose conformal confidence fan charts, CVaR risk slider, persona sensitivity radar, and automated C-Suite narrative generation.	tests/test_api_v2.py & Playwright E2E UI tests passing 100%.

Table 2: Structured engineering sprint plan for the implementation of NAOMI 2.0.

7. Business Value Proposition & Strategic Conclusion

By upgrading NAOMI from a standalone prototype to the **NAOMI 2.0 Unified Decision Companion**, enterprise retail leadership gains three unprecedented strategic capabilities:

1. **Zero Hallucinated Demand Increases:** By enforcing Simchi-Levi et al.'s shape constraints, NAOMI mathematically eliminates the dangerous upward-sloping demand curves that plague 60% of traditional ML retail models, safeguarding millions in operating margin.
2. **Distribution-Free Supply Chain Resilience:** By integrating CoRAL-RF's conformal prediction intervals, retail category managers can set safety stocks calibrated to exact empirical risk thresholds (e.g. 95% service level), cutting proxy inventory holding costs by 4.5% to 5.5%.
3. **Board-Level Risk Governance:** By combining Leeroy & Leeroy's counterfactual macroeconomic stress-testing with Columbia's CVaR optimization, C-suite executives can stress-test pricing policies against FX devaluation, stagflation, and supplier inflation with full visibility into worst-case tail risks.

Conclusion & Approval Recommendation: This research report presents a comprehensive, mathematically sound, and operationally viable blueprint. Upon executive approval of this proposal, the engineering team will immediately initiate Sprint 1 (feat/separable-elasticity) to deliver the next generation of AI-powered market intelligence.

8. Formal Academic References

[1] Zizi, M., Chafi, A., & El Hammoumi, M. (2026). A conformal LLM-guided random forest–LSTM forecaster for uncertainty-aware retail demand and price-scenario decision support. *Results in Engineering*, 32, 112272. Elsevier.

[2] Leeroy, W., & Leeroy, G. C. (2025). AI as a Decision Companion: Supporting Executive Pricing and FX Decisions in Global Enterprises Through LSTM Forecasting. *Journal of Risk and Financial Management*, 18(10), 542. MDPI.

[3] Li, M., Simchi-Levi, D., Tan, R., Wang, C., & Wu, M. X. (2024–2025). Contextual Offline Demand Learning and Pricing with Separable Models. *SSRN Working Paper No. 4619018*, MIT Institute for Data, Systems, and Society (IDSS).

[4] Huang, C., & Wang, K. (2026). LLM-Powered Virtual Population for Demand Simulation and Pricing. *SSRN Working Paper No. 6941578*, Department of Industrial Engineering and Operations Research, Columbia University.

[5] Vovk, V., Gammerman, A., & Shafer, G. (2005). *Algorithmic Learning in a Random World*. Springer Science & Business Media.

[6] Rockafellar, R. T., & Uryasev, S. (2000). Optimization of conditional value-at-risk. *Journal of Risk*, 2(3), 21–42.

[7] Simon, H. A. (1997). *Administrative Behavior: A Study of Decision-Making Processes in Administrative Organizations* (4th ed.). Free Press.

[8] Makridakis, S., Spiliotis, E., & Assimakopoulos, V. (2022). The M5 accuracy competition: Results, findings, and conclusions. *International Journal of Forecasting*, 38(4), 1346–1364.